



Weather + Lice Analysis

```
In [1]: from pathlib import Path
import os
import sys
import warnings

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tools.sm_exceptions import ConvergenceWarning

warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=ConvergenceWarning)

def find_project_root() -> Path:
    cwd = Path.cwd().resolve()
    for candidate in [cwd, cwd.parent]:
        if (candidate / "CA2").exists() and (candidate / "data").exists():
            return candidate
    return cwd

PROJECT_ROOT = find_project_root()
DATA_DIR = PROJECT_ROOT / "data"
CA2_DIR = PROJECT_ROOT / "CA2"
OUTPUT_DIR = CA2_DIR / "outputs"
OUTPUT_DIR.mkdir(exist_ok=True)

USE_CASSANDRA = 1
TARGET_YEAR = 2022

print(f"PROJECT_ROOT: {PROJECT_ROOT}")
print(f"USE_CASSANDRA: {USE_CASSANDRA}")
print(f"TARGET_YEAR: {TARGET_YEAR}")
```

```
PROJECT_ROOT: /Users/endreasgard/NMBU/IND320/IND320
USE_CASSANDRA: 1
TARGET_YEAR: 2022
```

1) Load locality dataset and create pivot analyses

```
In [2]: def classify_region(lat: float) -> str:
    if pd.isna(lat):
        return "Unknown"
```

```

    if lat < 63.81:
        return "South"
    if lat <= 68.85:
        return "Middle"
    return "North"

def load_locations() -> pd.DataFrame:
    candidates = [
        DATA_DIR / "locations_df.csv",
        DATA_DIR / "all.csv",
        PROJECT_ROOT / "CA1" / "fishhealth.csv",
    ]
    for candidate in candidates:
        if candidate.exists():
            print(f"Loading locality data from: {candidate}")
            return pd.read_csv(candidate)
    raise FileNotFoundError("No locality dataset found in expected paths.")

locations = load_locations().copy()

rename_map = {
    "localityNo": "localityno",
    "avgAdultFemaleLice": "avgadultfemalelice",
    "hasPd": "haspd",
    "hasIla": "hasila",
}
locations = locations.rename(columns=rename_map)

required = ["week", "avgadultfemalelice", "lat", "haspd", "hasila"]
missing = [c for c in required if c not in locations.columns]
if missing:
    raise ValueError(f"Missing locality columns: {missing}")

if "region" not in locations.columns:
    locations["region"] = locations["lat"].apply(classify_region)
else:
    locations["region"] = locations["lat"].apply(classify_region)

if "year" in locations.columns:
    year_filtered = locations[locations["year"] == TARGET_YEAR]
    if not year_filtered.empty:
        locations = year_filtered

week_lice = (
    locations.groupby("week", dropna=True)[ "avgadultfemalelice"]
    .mean()
    .dropna()
    .sort_index()
)
pd_ila_lat = (

```

```

        locations.pivot_table(
            index=["haspd", "hasila"],
            values="lat",
            aggfunc="mean",
        )
        .reset_index()
        .sort_values(["haspd", "hasila"])
    )

week_region_pivot = (
    locations.pivot_table(
        index="week",
        columns="region",
        values="avgadultfemalelice",
        aggfunc="mean",
    )
    .sort_index()
)

print(f"Locations rows used: {len(locations)}")
print("PD/ILA pivot sample:")
print(pd_ila_lat.head(10))
week_region_pivot.head(10)

```

Loading locality data from: /Users/endreasgard/NMBU/IND320/IND320/data/locations_df.csv

Locations rows used: 89,084

PD/ILA pivot sample:

	haspd	hasila	lat
0	False	False	64.088320
1	False	True	64.956952
2	True	False	61.868237
3	True	True	59.296028

Out[2]: **region** **Middle** **North** **South**

week			
1	0.140383	0.119362	0.178955
2	0.139032	0.115978	0.184441
3	0.147090	0.140115	0.188786
4	0.150056	0.147386	0.191480
5	0.166250	0.127595	0.200366
6	0.153314	0.136375	0.195725
7	0.163457	0.130303	0.191679
8	0.150616	0.158971	0.190577
9	0.148274	0.097931	0.190888
10	0.144562	0.145738	0.189809


```
In [3]: fig, axes = plt.subplots(1, 2, figsize=(15, 5))

for region in c for c in ["South", "Middle", "North"] if c in week_region_pivot:
    axes[0].plot(week_region_pivot.index, week_region_pivot[region], label=region)

axes[0].set_title("Average adult female lice by region")
axes[0].set_xlabel("Week")
axes[0].set_ylabel("Lice count")
axes[0].grid(alpha=0.3)
axes[0].legend()

heat_df = week_region_pivot.copy()
if not heat_df.empty:
    sns.heatmap(heat_df.T, cmap="YlGnBu", ax=axes[1])
    axes[1].set_title("Week-region heatmap (avg adult female lice)")
    axes[1].set_xlabel("Week")
    axes[1].set_ylabel("Region")

plt.tight_layout()
plt.show()
```



2) Load synchronized weather + lice features (single locality)

```
In [4]: def load_weather_lice_features() -> pd.DataFrame:
    preferred = DATA_DIR / "sar_weather_lice.csv"
    if preferred.exists():
        print(f"Loading merged feature data from: {preferred}")
        return pd.read_csv(preferred)

    lice_candidates = [DATA_DIR / "lice_df_2.csv", DATA_DIR / "lice.csv"]
    weather_candidates = [DATA_DIR / "weather_lice_lag.csv", DATA_DIR / "weat
```

```

        lice_df = pd.read_csv(candidate)
        break

    for candidate in weather_candidates:
        if candidate.exists():
            weather_df = pd.read_csv(candidate)
            break

    if lice_df is None or weather_df is None:
        raise FileNotFoundError("Could not build feature dataset from local file")

    lice_df = lice_df.rename(columns={
        "avgAdultFemaleLice": "avgadultfemalelice",
        "avgMobileLice": "avgmobilelice",
        "avgStationaryLice": "avgstationarylice",
        "seaTemperature": "seatemperature",
    })

    weather_df = weather_df.rename(columns={
        "mean(air_temperature P1D)": "mean_air_temperature",
        "mean(relative_humidity P1D)": "mean_relative_humidity",
        "mean(wind_speed P1D)": "mean_wind_speed",
        "sum(precipitation_amount P1D)": "sum_precipitation_amount",
    })

    merge_keys = [k for k in ["year", "week"] if k in lice_df.columns and k in weather_df.columns]
    if not merge_keys and "week" in lice_df.columns and "week" in weather_df.columns:
        merge_keys = ["week"]

    df = pd.merge(lice_df, weather_df, how="inner", on=merge_keys)

    if "smoothed_avgadultfemalelice" not in df.columns and "avgadultfemalelice" in df.columns:
        df["smoothed_avgadultfemalelice"] = df["avgadultfemalelice"].rolling(window=2).mean()
    if "smoothed_avgmobilelice" not in df.columns and "avgmobilelice" in df.columns:
        df["smoothed_avgmobilelice"] = df["avgmobilelice"].rolling(window=2).mean()
    if "smoothed_avgstationarylice" not in df.columns and "avgstationarylice" in df.columns:
        df["smoothed_avgstationarylice"] = df["avgstationarylice"].rolling(window=2).mean()

    for base_col in ["mean_air_temperature", "mean_relative_humidity", "mean_wind_speed"]:
        lag_col = f"{base_col}_lag1"
        if base_col in df.columns and lag_col not in df.columns:
            df[lag_col] = df[base_col].shift(1)

    return df

feature_df = load_weather_lice_features().copy()
feature_df = feature_df.sort_values("week").reset_index(drop=True)

core_cols = [
    "week",
    "seatemperature",
    "avgadultfemalelice",
]

```

```

    "avgmobilelice",
    "avgstationarylice",
    "mean_air_temperature",
    "mean_relative_humidity",
    "mean_wind_speed",
    "sum_precipitation_amount",
    "mean_air_temperature_lag1",
    "mean_relative_humidity_lag1",
    "mean_wind_speed_lag1",
    "sum_precipitation_amount_lag1",
]

missing_core = [c for c in ["week", "seatemperature", "avgadultfemalelice", "a
if missing_core:
    raise ValueError(f"Missing required feature columns: {missing_core}")

for col in core_cols:
    if col in feature_df.columns:
        feature_df[col] = pd.to_numeric(feature_df[col], errors="coerce")

feature_df = feature_df.dropna(subset=["week", "seatemperature", "avgadultfema
print(f"Feature rows available: {len(feature_df)}")
feature_df.head(10)

```

Loading merged feature data from: /Users/endreasgard/NMBU/IND320/IND320/data/sa
r_weather_lice.csv
Feature rows available: 26

Out[4]:

	week	avgadultfemalelice	avgmobilelice	avgstationarylice	mean_air_temper
0	3	0.04	0.29	0.0	3.4!
1	4	0.05	0.27	0.0	1.9!
2	5	0.07	0.27	0.0	-1.8!
3	6	0.06	0.23	0.0	1.4!
4	7	0.09	0.42	0.0	-1.7!
5	9	0.12	0.55	0.0	3.4!
6	11	0.07	0.38	0.0	6.5!
7	13	0.19	0.47	0.0	0.5!
8	15	0.12	0.58	0.0	5.4!
9	16	0.16	0.69	0.0	8.8!

In [5]:

```

corr_cols = [
    "seatemperature",
    "avgadultfemalelice",
    "avgmobilelice",
    "avgstationarylice",
    "mean_air_temperature",

```


```

    "mean_relative_humidity",
    "mean_wind_speed",
    "sum_precipitation_amount",
]

corr_cols = [c for c in corr_cols if c in feature_df.columns]
corr = feature_df[corr_cols].corr(numeric_only=True)

plt.figure(figsize=(9, 7))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", square=True)
plt.title("Correlation matrix (weather and lice features)")
plt.tight_layout()
plt.show()

```



3) Sliding Window Correlation (sea temperature vs lice)

```
In [6]: def sliding_window_correlation(temp: pd.Series, lice: pd.Series, window: int,
```

```

"""Positive lag means sea temperature leads lice by `lag` weeks."""
temp = temp.reset_index(drop=True)
lice = lice.reset_index(drop=True)

if lag > 0:
    temp_aligned = temp.iloc[: -lag].reset_index(drop=True)
    lice_aligned = lice.iloc[lag : ].reset_index(drop=True)
elif lag < 0:
    temp_aligned = temp.iloc[-lag : ].reset_index(drop=True)
    lice_aligned = lice.iloc[: lag].reset_index(drop=True)
else:
    temp_aligned = temp
    lice_aligned = lice

max_start = len(temp_aligned) - window
if max_start < 0:
    return pd.Series(dtype=float)

corr_values = []
window_index = []

for start in range(max_start + 1):
    temp_slice = temp_aligned.iloc[start : start + window]
    lice_slice = lice_aligned.iloc[start : start + window]
    corr_values.append(temp_slice.corr(lice_slice))
    window_index.append(start + 1)

return pd.Series(corr_values, index=window_index)

def find_best_window_lag(temp: pd.Series, lice: pd.Series, windows, lags):
    best = None
    for window in windows:
        for lag in lags:
            corr_series = sliding_window_correlation(temp, lice, window=window)
            if corr_series.empty:
                continue
            score = corr_series.abs().mean()
            if best is None or score > best["score"]:
                best = {
                    "window": window,
                    "lag": lag,
                    "score": float(score),
                    "series": corr_series,
                }
    return best

temp_series = feature_df["seatemperature"].astype(float)
lice_series_map = {
    "avgadultfemalelice": feature_df["avgadultfemalelice"].astype(float),
    "avgmobilelice": feature_df["avgmobilelice"].astype(float),
    "avgstationarylice": feature_df["avgstationarylice"].astype(float),
}

```

```

}

windows = [4, 6, 8, 10]
lags = [0, 1, 2, 3]

best_configs = {}
for key, lice_series in lice_series_map.items():
    best = find_best_window_lag(temp_series, lice_series, windows=windows, lag=lags)
    best_configs[key] = best

best_summary = pd.DataFrame(
    [
        {
            "series": key,
            "best_window": cfg["window"],
            "best_lag": cfg["lag"],
            "mean_abs_correlation": cfg["score"],
        }
        for key, cfg in best_configs.items()
        if cfg is not None
    ]
)
best_summary

```

```

/Users/endreasgard/miniconda3/envs/personlig311/lib/python3.11/site-packages/numpy/lib/_function_base_impl.py:3065: RuntimeWarning: invalid value encountered
in divide
    c /= stddev[:, None]
/Users/endreasgard/miniconda3/envs/personlig311/lib/python3.11/site-packages/numpy/lib/_function_base_impl.py:3066: RuntimeWarning: invalid value encountered
in divide
    c /= stddev[None, :]

```

Out[6]:

	series	best_window	best_lag	mean_abs_correlation
0	avgadultfemalelice	4	0	0.610241
1	avgmobilelice	4	0	0.606241
2	avgstationarylice	6	3	0.774091

In [7]:

```

fig, axes = plt.subplots(3, 1, figsize=(12, 12), sharex=False)

for ax, (series_name, cfg) in zip(axes, best_configs.items()):
    if cfg is None:
        ax.set_title(f"{series_name}: no valid window/lag config")
        continue

    corr_series = cfg["series"]
    ax.plot(corr_series.index, corr_series.values, marker="o")
    ax.axhline(0, color="black", linewidth=1)
    ax.set_title(
        f"{series_name} | best window={cfg['window']}, lag={cfg['lag']} | "
    )


```

```

        f"mean |corr|={cfg['score']:.3f}"
    )
ax.set_xlabel("Window start index")
ax.set_ylabel("Correlation")
ax.grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



4) ARIMAX forecasting for sea temperature

Model A: weather variables only. Model B: weather + lagged weather variables.

Model C: Model B minus the least influential variable.

```
In [8]: def fit_sarimax_with_holdout(df: pd.DataFrame, exog_cols, label: str):
    model_df = df[["week", "seatemperature"] + exog_cols].dropna().reset_index
    if len(model_df) < 14:
        raise ValueError(f"Not enough rows to train/test for {label}.")
    y = model_df["seatemperature"].astype(float)
    X = model_df[exog_cols].astype(float)

    split_idx = max(12, len(model_df) - 6)
    y_train, y_test = y.iloc[:split_idx], y.iloc[split_idx:]
    X_train, X_test = X.iloc[:split_idx], X.iloc[split_idx:]

    model = SARIMAX(
        y_train,
        exog=X_train,
        order=(1, 0, 0),
        trend="c",
        enforce_stationarity=False,
        enforce_invertibility=False,
    )
    result = model.fit(disp=False, maxiter=200)

    forecast = result.get_forecast(steps=len(y_test), exog=X_test).predicted
    rmse = float(np.sqrt(np.mean((y_test.values - forecast.values) ** 2)))

    output = {
        "label": label,
        "result": result,
        "aic": float(result.aic),
        "rmse": rmse,
        "weeks_test": model_df["week"].iloc[split_idx:].to_numpy(),
        "actual_test": y_test.to_numpy(),
        "pred_test": forecast.to_numpy(),
        "exog_cols": list(exog_cols),
    }
    return output

base_cols = [
    "mean_air_temperature",
    "mean_relative_humidity",
    "mean_wind_speed",
    "sum_precipitation_amount",
]
lag_cols = [
    "mean_air_temperature_lag1",
    "mean_relative_humidity_lag1",
    "mean_wind_speed_lag1",
    "sum_precipitation_amount_lag1",
]

base_cols = [c for c in base_cols if c in feature_df.columns]
lag_cols = [c for c in lag_cols if c in feature_df.columns]
```

```

model_a = fit_sarimax_with_holdout(feature_df, base_cols, "A_base_weather")
model_b = fit_sarimax_with_holdout(feature_df, base_cols + lag_cols, "B_base_plus_lagged")

params_b = model_b["result"].params.drop(labels=["sigma2"], errors="ignore")
coef_b = params_b[[c for c in params_b.index if c in model_b["exog_cols"]]]
least_influential = coef_b.abs().sort_values().index[0]

reduced_cols = [c for c in model_b["exog_cols"] if c != least_influential]
model_c = fit_sarimax_with_holdout(feature_df, reduced_cols, "C_reduced")

metrics = pd.DataFrame([
    {"model": model_a["label"], "aic": model_a["aic"], "rmse_test": model_a["rmse_test"]},
    {"model": model_b["label"], "aic": model_b["aic"], "rmse_test": model_b["rmse_test"]},
    {"model": model_c["label"], "aic": model_c["aic"], "rmse_test": model_c["rmse_test"]}
]).sort_values("rmse_test")

```

metrics

Out[8]:

	model	aic	rmse_test	n_exog	removed_variable
0	A_base_weather	91.542499	1.145872	4	-
2	C_reduced	71.647437	2.194353	7	mean_relative_humidity_lag1
1	B_base_plus_lagged	73.636098	2.233651	8	-

In [9]:


```

plt.figure(figsize=(12, 5))

for model in [model_a, model_b, model_c]:
    plt.plot(model["weeks_test"], model["pred_test"], marker="o", label=f"{model['label']}")

plt.plot(model_a["weeks_test"], model_a["actual_test"], color="black", linewidth=2)
plt.title("Test-set forecast comparison")
plt.xlabel("Week")
plt.ylabel("Sea temperature")
plt.grid(alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()

```



5) Cassandra export

```
In [10]: if USE_CASSANDRA:
    try:
        from cassandra.cluster import Cluster

        cluster = Cluster(["localhost"], port=9042)
        session = cluster.connect()
        session.execute(
            """
            CREATE KEYSPACE IF NOT EXISTS fishhealth
            WITH REPLICATION = {'class': 'SimpleStrategy', 'replication_factor': 1}
            """
        )
        session.set_keyspace("fishhealth")

        session.execute(
            """
            CREATE TABLE IF NOT EXISTS weather_lice_cv (
                week int,
                seatemperature double,
                avgadultfemalelice double,
                avgmobilelice double,
                avgstationarylice double,
                mean_air_temperature double,
                mean_relative_humidity double,
                mean_wind_speed double,
                sum_precipitation_amount double,
                PRIMARY KEY (week)
            )
            """
        )

        cols = [
            "week",
            "seatemperature",

```

```

        "avgadultfemalelice",
        "avgmobilelice",
        "avgstationarylice",
        "mean_air_temperature",
        "mean_relative_humidity",
        "mean_wind_speed",
        "sum_precipitation_amount",
    ]

    insert_stmt = session.prepare(
        """
        INSERT INTO weather_lice_cv
        (week, seatemperature, avgadultfemalelice, avgmobilelice, avgstationarytemperature,
         mean_air_temperature, mean_relative_humidity, mean_wind_speed, sum_precipitation_amount)
        VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
        """
    )

    inserted = 0
    for row in feature_df[cols].dropna(subset=["week"]).itertuples(index=False):
        session.execute(
            insert_stmt,
            (
                int(row.week),
                float(row.seatemperature) if pd.notna(row.seatemperature) else None,
                float(row.avgadultfemalelice) if pd.notna(row.avgadultfemalelice) else None,
                float(row.avgmobilelice) if pd.notna(row.avgmobilelice) else None,
                float(row.avgstationarylice) if pd.notna(row.avgstationarytemperature) else None,
                float(row.mean_air_temperature) if pd.notna(row.mean_air_temperature) else None,
                float(row.mean_relative_humidity) if pd.notna(row.mean_relative_humidity) else None,
                float(row.mean_wind_speed) if pd.notna(row.mean_wind_speed) else None,
                float(row.sum_precipitation_amount) if pd.notna(row.sum_precipitation_amount) else None
            ),
        )
        inserted += 1

    print(f"Inserted {inserted} rows into fishhealth.weather_lice_cv")

except Exception as exc:
    print(f"Cassandra step skipped due to error: {exc}")
finally:
    if "cluster" in locals():
        cluster.shutdown()
else:
    print("Skipping Cassandra step. Set BW_USE_CASSANDRA=1 to enable.")

```

Inserted 26 rows into fishhealth.weather_lice_cv

```
In [11]: week_region_out = OUTPUT_DIR / "week_region_pivot.csv"
pd_ilat_out = OUTPUT_DIR / "pd_ilat_lat_pivot.csv"
features_out = OUTPUT_DIR / "weather_lice_features.csv"
metrics_out = OUTPUT_DIR / "forecast_metrics.csv"
sliding_out = OUTPUT_DIR / "sliding_window_best_configs.csv"
```

```
week_region_pivot.reset_index().to_csv(week_region_out, index=False)
pd_ila_lat.to_csv(pd_ila_out, index=False)
feature_df.to_csv(features_out, index=False)
metrics.to_csv(metrics_out, index=False)
best_summary.to_csv(sliding_out, index=False)

print(f"Saved: {week_region_out}")
print(f"Saved: {pd_ila_out}")
print(f"Saved: {features_out}")
print(f"Saved: {metrics_out}")
print(f"Saved: {sliding_out}")
```

```
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA2/outputs/week_region_pivot.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA2/outputs/pd_ila_lat_pivot.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA2/outputs/weather_lice_features.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA2/outputs/forecast_metrics.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA2/outputs/sliding_window_best_configs.csv
```