



CA3 Integrated

```
In [1]: import sys, pyspark  
print(sys.executable)  
print(pyspark.__version__)  
  
/Users/endreasgard/miniconda3/envs/personlig311/bin/python  
3.5.1
```

```
In [1]: from pathlib import Path  
import sys  
import warnings  
import subprocess  
  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
from sklearn.ensemble import GradientBoostingRegressor  
from sklearn.metrics import mean_absolute_error, mean_squared_error  
from sklearn.model_selection import GridSearchCV, TimeSeriesSplit  
from statsmodels.tools.sm_exceptions import ConvergenceWarning  
  
warnings.filterwarnings("ignore", category=FutureWarning)  
warnings.filterwarnings("ignore", category=ConvergenceWarning)  
  
  
def locate_root() -> Path:  
    cwd = Path.cwd().resolve()  
    candidates = [cwd, cwd.parent]  
    for candidate in candidates:  
        if (candidate / "common").exists() and (candidate / "data").exists():  
            return candidate  
    return cwd  
  
  
PROJECT_ROOT = locate_root()  
if str(PROJECT_ROOT) not in sys.path:  
    sys.path.insert(0, str(PROJECT_ROOT))  
  
from common.ca_shared import (  
    add_region_column,  
    best_window_lag,  
    create_cassandra_connection,  
    create_spark_session,  
    ensure_keyspace,  
    fit_sarimax_holdout,  
    load_settings,  
    read_cassandra_table_with_spark,  
    write_spark_df_to_cassandra,  
)
```

```
DATA_DIR = PROJECT_ROOT / "data"
OUTPUT_DIR = PROJECT_ROOT / "CA3" / "outputs"
OUTPUT_DIR.mkdir(exist_ok=True)

java_version = subprocess.check_output(["java", "-version"], stderr=subprocess.STDOUT)
# print("Java:", java_version)
# print("Project root:", PROJECT_ROOT)
```

```
In [2]: settings = load_settings()
TARGET_YEAR = settings["general"]["target_year"]
KEYSPACE = settings["cassandra"]["keyspace"]

cluster, session = create_cassandra_connection(settings)
ensure_keyspace(session, KEYSPACE)

try:
    spark = create_spark_session(settings, app_name="CA3CassandraSpark")
except Exception as exc:
    cluster.shutdown()
    raise RuntimeError(
        "Spark failed to start. Check version compatibility first: "
        "pyspark must be <=3.5.x for Cassandra connector support, and Java shd"
        f"Root cause: {exc}"
    ) from exc

print("Connected Cassandra keyspace:", KEYSPACE)
print("Spark version:", spark.version)

:: loading settings :: url = jar:file:/Users/endreasgard/miniconda3/envs/person
lig311/lib/python3.11/site-packages/pyspark/jars/ivy-2.5.1.jar!/org/apache/ivy/
core/settings/ivysettings.xml
```

```

Ivy Default Cache set to: /Users/endreasgard/.ivy2/cache
The jars for the packages stored in: /Users/endreasgard/.ivy2/jars
com.datastax.spark#spark-cassandra-connector_2.12 added as a dependency
:: resolving dependencies :: org.apache.spark#spark-submit-parent-f026f071-a45
a-487f-aa78-d8b32800e767;1.0
    confs: [default]
        found com.datastax.spark#spark-cassandra-connector_2.12;3.5.1 in centra
l
        found com.datastax.spark#spark-cassandra-connector-driver_2.12;3.5.1 in central
        found org.scala-lang.modules#scala-collection-compat_2.12;2.11.0 in central
        found org.apache.cassandra#java-driver-core-shaded;4.18.1 in central
        found com.datastax.oss#native-protocol;1.5.1 in central
        found com.datastax.oss#java-driver-shaded-guava;25.1-jre-graal-sub-1 in central
        found com.typesafe#config;1.4.1 in central
        found org.slf4j#slf4j-api;1.7.26 in central
        found io.dropwizard.metrics#metrics-core;4.1.18 in central
        found org.hdrhistogram#HdrHistogram;2.1.12 in central
        found org.reactivestreams#reactive-streams;1.0.3 in central
        found org.apache.cassandra#java-driver-mapper-runtime;4.18.1 in central
        found org.apache.cassandra#java-driver-query-builder;4.18.1 in central
        found org.apache.commons#commons-lang3;3.10 in central
        found com.thoughtworks.paranamer#paranamer;2.8 in central
        found org.scala-lang#scala-reflect;2.12.19 in central
:: resolution report :: resolve 205ms :: artifacts dl 6ms
    :: modules in use:
        com.datastax.oss#java-driver-shaded-guava;25.1-jre-graal-sub-1 from cen
tral in [default]
        com.datastax.oss#native-protocol;1.5.1 from central in [default]
        com.datastax.spark#spark-cassandra-connector-driver_2.12;3.5.1 from cen
tral in [default]
        com.datastax.spark#spark-cassandra-connector_2.12;3.5.1 from central in
[default]
        com.thoughtworks.paranamer#paranamer;2.8 from central in [default]
        com.typesafe#config;1.4.1 from central in [default]
        io.dropwizard.metrics#metrics-core;4.1.18 from central in [default]
        org.apache.cassandra#java-driver-core-shaded;4.18.1 from central in [de
fault]
        org.apache.cassandra#java-driver-mapper-runtime;4.18.1 from central in
[default]
        org.apache.cassandra#java-driver-query-builder;4.18.1 from central in
[default]
        org.apache.commons#commons-lang3;3.10 from central in [default]
        org.hdrhistogram#HdrHistogram;2.1.12 from central in [default]
        org.reactivestreams#reactive-streams;1.0.3 from central in [default]
        org.scala-lang#scala-reflect;2.12.19 from central in [default]
        org.scala-lang.modules#scala-collection-compat_2.12;2.11.0 from central
in [default]
        org.slf4j#slf4j-api;1.7.26 from central in [default]
-----
|           |           modules          ||   artifacts   |
|       conf    | number| search|dwnlded|evicted|| number|dwnlded|

```

```
|      default      |    16 |    0 |    0 |    0 ||    16 |    0 |
-----
```

:: retrieving :: org.apache.spark#spark-submit-parent-f026f071-a45a-487f-aa78-d8b32800e767
conf: [default]
0 artifacts copied, 16 already retrieved (0kB/5ms)
26/02/21 14:43:16 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Connected Cassandra keyspace: fishhealth
Spark version: 3.5.1

1) Prepare Cassandra tables from local project data

```
In [4]: session.execute(
    """
        CREATE TABLE IF NOT EXISTS ca3_localities (
            week int,
            localityno int,
            year int,
            avgadultfemalelice double,
            haspd boolean,
            hasila boolean,
            lat double,
            lon double,
            region text,
            PRIMARY KEY ((week), localityno, year)
        )
    """
)

session.execute(
    """
        CREATE TABLE IF NOT EXISTS ca3_weather_lice (
            week int,
            seatemperature double,
            avgadultfemalelice double,
            avgmobilelice double,
            avgstationarylice double,
            mean_air_temperature double,
            mean_relative_humidity double,
            mean_wind_speed double,
            sum_precipitation_amount double,
            mean_air_temperature_lag1 double,
            mean_relative_humidity_lag1 double,
            mean_wind_speed_lag1 double,
            sum_precipitation_amount_lag1 double,
    """
)
```

```

        PRIMARY KEY (week)
    )
"""

)

loc_count = session.execute("SELECT count(*) FROM ca3_localities").one().count
feat_count = session.execute("SELECT count(*) FROM ca3_weather_lice").one().count
print("Existing rows in ca3_localities:", loc_count)
print("Existing rows in ca3_weather_lice:", feat_count)

if loc_count == 0:
    locations_path = DATA_DIR / "locations_df.csv"
    if not locations_path.exists():
        raise FileNotFoundError(f"Missing file: {locations_path}")

    loc_pd = pd.read_csv(locations_path)
    loc_pd = loc_pd.rename(
        columns={
            "localityNo": "localityno",
            "avgAdultFemaleLice": "avgadultfemalelice",
            "hasPd": "haspd",
            "hasIla": "hasila",
        }
    )

loc_pd = add_region_column(loc_pd, lat_col="lat", target_col="region")

if "year" in loc_pd.columns:
    filtered = loc_pd[loc_pd["year"] == TARGET_YEAR]
    if not filtered.empty:
        loc_pd = filtered

loc_pd = loc_pd[[
    "week", "localityno", "year", "avgadultfemalelice",
    "haspd", "hasila", "lat", "lon", "region"
]].copy()

loc_pd["week"] = pd.to_numeric(loc_pd["week"], errors="coerce").astype("Int64")
loc_pd["localityno"] = pd.to_numeric(loc_pd["localityno"], errors="coerce")
loc_pd["year"] = pd.to_numeric(loc_pd["year"], errors="coerce").astype("Int64")

loc_pd = loc_pd.dropna(subset=["week", "localityno", "year", "lat", "lon"])
loc_pd["week"] = loc_pd["week"].astype(int)
loc_pd["localityno"] = loc_pd["localityno"].astype(int)
loc_pd["year"] = loc_pd["year"].astype(int)

loc_spark = spark.createDataFrame(loc_pd)
write_spark_df_to_cassandra(loc_spark, KEYSPACE, "ca3_localities", mode="append")

loc_count = session.execute("SELECT count(*) FROM ca3_localities").one().count
print("Rows inserted into ca3_localities:", loc_count)

if feat_count == 0:

```

```

feature_path = DATA_DIR / "sar_weather_lice.csv"
if not feature_path.exists():
    raise FileNotFoundError(f"Missing file: {feature_path}")

feat_pd = pd.read_csv(feature_path)

needed = [
    "week",
    "seatemperature",
    "avgadultfemalelice",
    "avgmobilelice",
    "avgstationarylice",
    "mean_air_temperature",
    "mean_relative_humidity",
    "mean_wind_speed",
    "sum_precipitation_amount",
    "mean_air_temperature_lag1",
    "mean_relative_humidity_lag1",
    "mean_wind_speed_lag1",
    "sum_precipitation_amount_lag1",
]
missing = [c for c in needed if c not in feat_pd.columns]
if missing:
    raise ValueError(f"Missing columns in sar_weather_lice.csv: {missing}")

feat_pd = feat_pd[needed].copy()
feat_pd["week"] = pd.to_numeric(feat_pd["week"], errors="coerce").astype("int")
feat_pd = feat_pd.dropna(subset=["week"])
feat_pd["week"] = feat_pd["week"].astype(int)

feat_spark = spark.createDataFrame(feat_pd)
write_spark_df_to_cassandra(feat_spark, KEYSPACE, "ca3_weather_lice", mode="append")

feat_count = session.execute("SELECT count(*) FROM ca3_weather_lice").one()
print("Rows inserted into ca3_weather_lice:", feat_count)

```

Existing rows in ca3_localities: 89084
 Existing rows in ca3_weather_lice: 26

2) Read data back through Spark from Cassandra

```

In [5]: locations_spark = read_cassandra_table_with_spark(spark, KEYSPACE, "ca3_localities")
features_spark = read_cassandra_table_with_spark(spark, KEYSPACE, "ca3_weather_lice")

locations_df = locations_spark.toPandas().sort_values(["week", "localityno"])
feature_df = features_spark.toPandas().sort_values("week")

print("locations_df shape:", locations_df.shape)
print("feature_df shape:", feature_df.shape)
locations_df.head(5)

```

```

locations_df shape: (89084, 9)
feature_df shape: (26, 13)

Out[5]:
```

	week	localityno	year	avgadultfemalelice	hasila	haspd		lat	
32553	1	10029	2022		NaN	False	False	59.371235	5.2
32554	1	10037	2022		NaN	False	False	60.339880	4.9
32555	1	10040	2022		NaN	False	False	59.880650	5.1
32556	1	10041	2022	0.38	False	False	60.054317	5.0	
32557	1	10044	2022		NaN	False	False	59.862167	5.1

3) Pivoting and regional analysis

```

In [6]: weekly_avg = (
    locations_df.groupby("week", dropna=True)[ "avgadultfemalelice"]
    .mean()
    .reset_index(name="avgadultfemalelice")
)

pivot_table_pd_ila = (
    locations_df.pivot_table(
        index=[ "haspd", "hasila"],
        values="lat",
        aggfunc="mean",
    )
    .reset_index()
)

pivot_table_week_region = (
    locations_df.pivot_table(
        index="week",
        columns="region",
        values="avgadultfemalelice",
        aggfunc="mean",
    )
    .sort_index()
)

fig, axes = plt.subplots(2, 2, figsize=(16, 10))

axes[0, 0].plot(weekly_avg["week"], weekly_avg["avgadultfemalelice"], marker="o")
axes[0, 0].set_title("Weekly average adult female lice")
axes[0, 0].set_xlabel("Week")
axes[0, 0].set_ylabel("Lice count")
axes[0, 0].grid(alpha=0.3)

for region in [c for c in ["South", "Middle", "North"] if c in pivot_table_week_region]:
    axes[0, 1].plot(pivot_table_week_region.index, pivot_table_week_region[region])

```

```

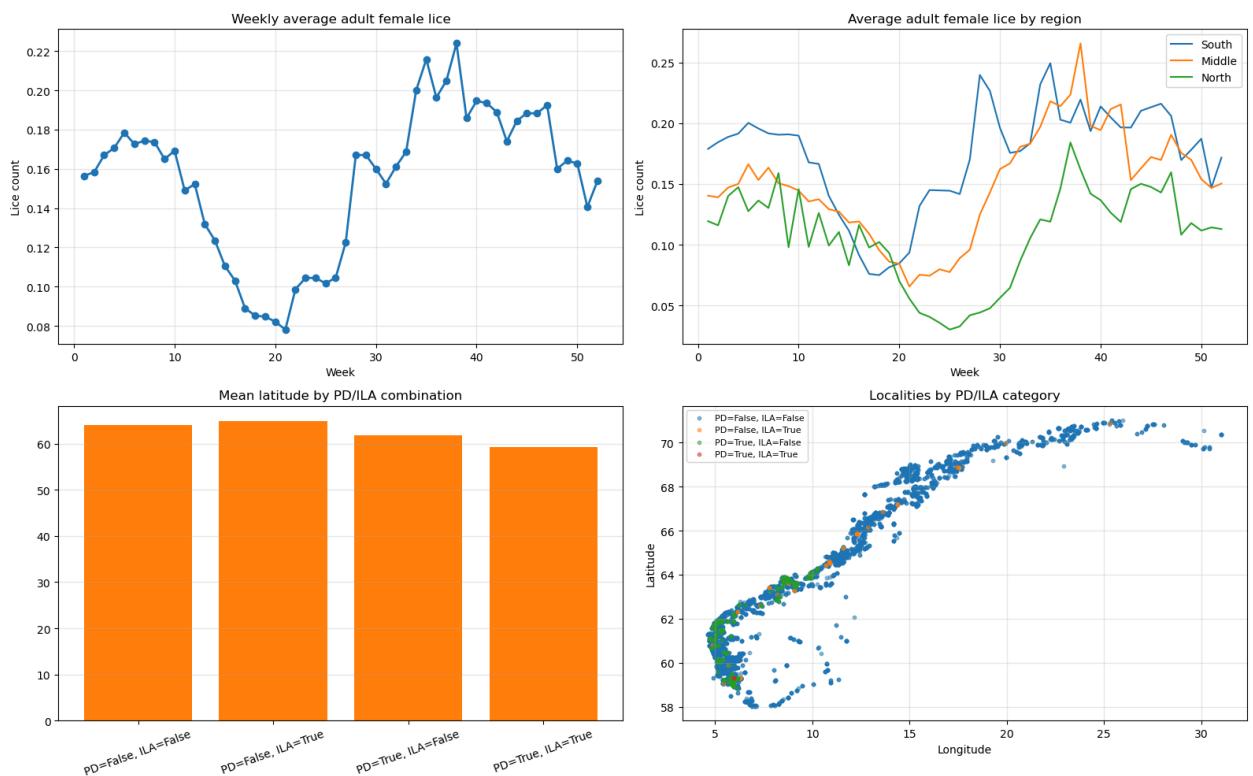
axes[0, 1].set_title("Average adult female lice by region")
axes[0, 1].set_xlabel("Week")
axes[0, 1].set_ylabel("Lice count")
axes[0, 1].legend()
axes[0, 1].grid(alpha=0.3)

tmp = pivot_table_pd_ilा.copy()
tmp["label"] = "PD=" + tmp["haspd"].astype(str) + ", ILA=" + tmp["hasila"].astype(str)
axes[1, 0].bar(tmp["label"], tmp["lat"], color="tab:orange")
axes[1, 0].set_title("Mean latitude by PD/ILA combination")
axes[1, 0].tick_params(axis="x", rotation=20)
axes[1, 0].grid(axis="y", alpha=0.3)

sample_map = locations_df[["lon", "lat", "haspd", "hasila"]].dropna().copy()
if len(sample_map) > 5000:
    sample_map = sample_map.sample(5000, random_state=42)
sample_map["category"] = "PD=" + sample_map["haspd"].astype(str) + ", ILA=" +
for category, group in sample_map.groupby("category"):
    axes[1, 1].scatter(group["lon"], group["lat"], s=10, alpha=0.5, label=category)
axes[1, 1].set_title("Localities by PD/ILA category")
axes[1, 1].set_xlabel("Longitude")
axes[1, 1].set_ylabel("Latitude")
axes[1, 1].legend(fontsize=8)
axes[1, 1].grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



4) Sliding-window correlations (sea temperature vs lice)

```
In [7]: temp = feature_df["seatemperature"].astype(float)
series_map = {
    "avgadultfemalelice": feature_df["avgadultfemalelice"].astype(float),
    "avgmobilelice": feature_df["avgmobilelice"].astype(float),
    "avgstationarylice": feature_df["avgstationarylice"].astype(float),
}

windows = [4, 6, 8, 10]
lags = [0, 1, 2, 3]

best_cfg = {name: best_window_lag(temp, series, windows, lags) for name, series
            in series_map.items()}

best_cfg_df = pd.DataFrame([
    {
        "series": name,
        "best_window": cfg["window"],
        "best_lag": cfg["lag"],
        "mean_abs_correlation": cfg["score"],
    }
    for name, cfg in best_cfg.items()
    if cfg is not None
])

best_cfg_df
```

```
/Users/endreasgard/miniconda3/envs/personlig311/lib/python3.11/site-packages/numpy/lib/_function_base_impl.py:3065: RuntimeWarning: invalid value encountered
in divide
    c /= stddev[:, None]
/Users/endreasgard/miniconda3/envs/personlig311/lib/python3.11/site-packages/numpy/lib/_function_base_impl.py:3066: RuntimeWarning: invalid value encountered
in divide
    c /= stddev[None, :]
```

```
Out[7]:
```

	series	best_window	best_lag	mean_abs_correlation
0	avgadultfemalelice	4	0	0.610241
1	avgmobilelice	4	0	0.606241
2	avgstationarylice	6	3	0.774091

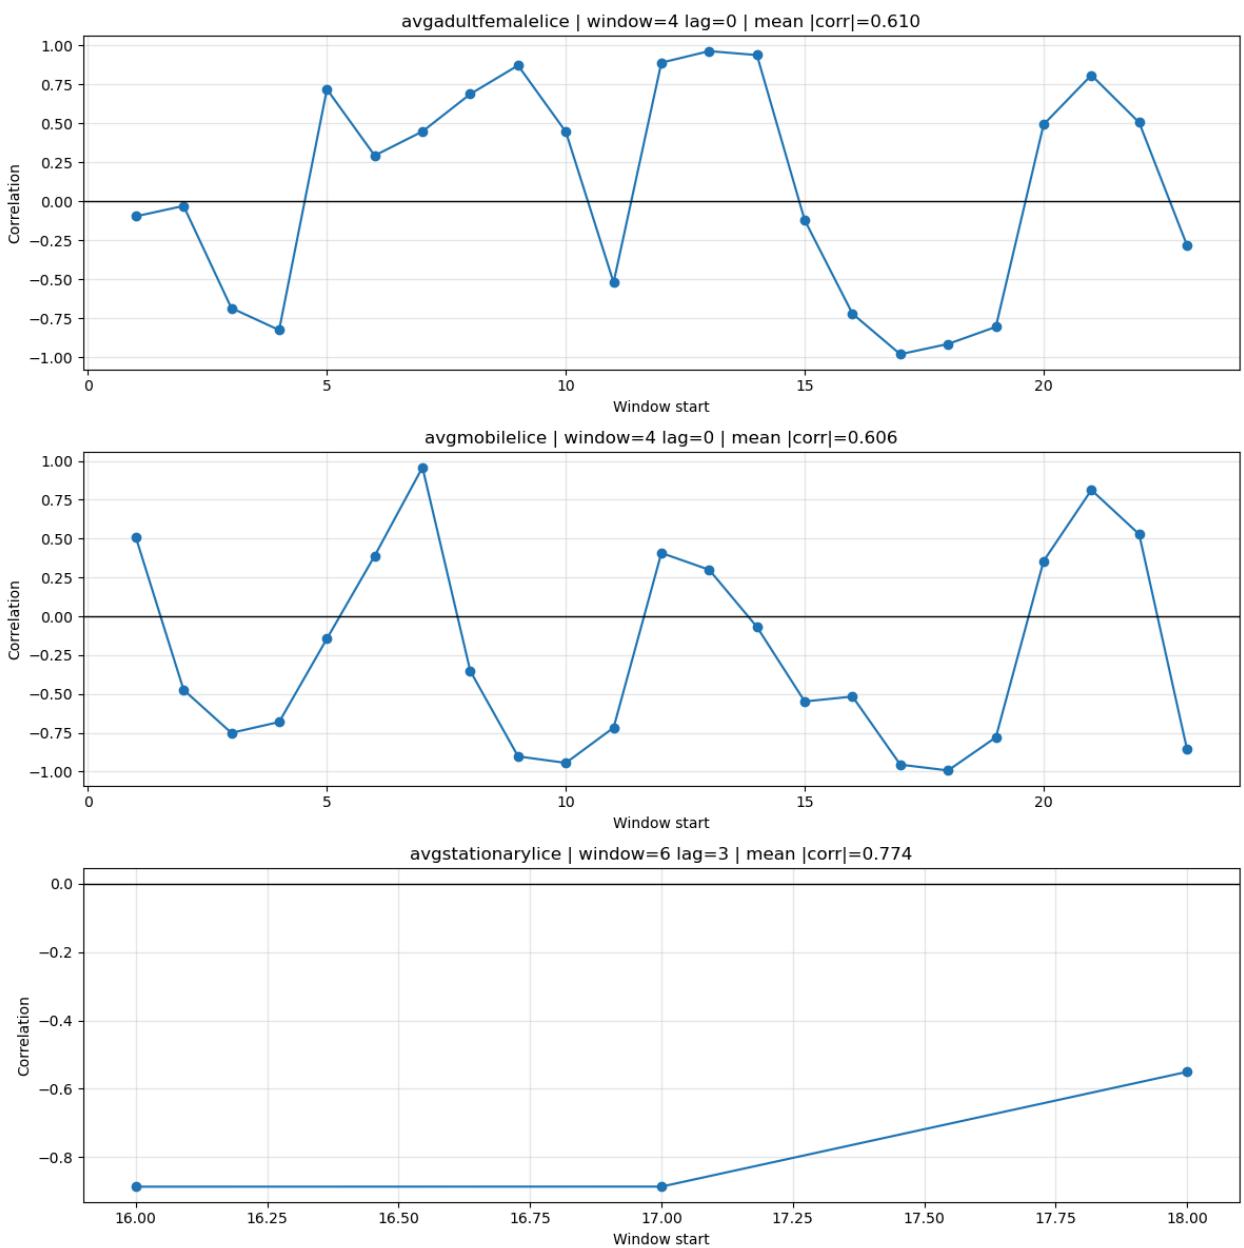
```
In [8]: fig, axes = plt.subplots(3, 1, figsize=(12, 12))
for ax, (name, cfg) in zip(axes, best_cfg.items()):
    if cfg is None:
        ax.set_title(f"{name}: no valid config")
        continue
    s = cfg["series"]
    ax.plot(s.index, s.values, marker="o")
```

```

    ax.axhline(0, color="black", linewidth=1)
    ax.set_title(f"{name} | window={cfg['window']} | lag={cfg['lag']} | mean |corr|={corr:.4f}")
    ax.set_xlabel("Window start")
    ax.set_ylabel("Correlation")
    ax.grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



5) Forecasting (ARIMAX + ML baseline)

```
In [9]: base_cols = [
    "mean_air_temperature",
    "mean_relative_humidity",
```

```

        "mean_wind_speed",
        "sum_precipitation_amount",
    ]
lag_cols = [
    "mean_air_temperature_lag1",
    "mean_relative_humidity_lag1",
    "mean_wind_speed_lag1",
    "sum_precipitation_amount_lag1",
]

base_cols = [c for c in base_cols if c in feature_df.columns]
lag_cols = [c for c in lag_cols if c in feature_df.columns]

arimax_a = fit_sarimax_holdout(feature_df, "seatemperature", base_cols, "ARIMA_A")
arimax_b = fit_sarimax_holdout(feature_df, "seatemperature", base_cols + lag_cols, "ARIMA_B")

params_b = arimax_b["result"].params.drop(labels=["sigma2"], errors="ignore")
coef_b = params_b[[c for c in params_b.index if c in (base_cols + lag_cols)]]
least_influential = coef_b.abs().sort_values().index[0]

reduced_cols = [c for c in (base_cols + lag_cols) if c != least_influential]
arimax_c = fit_sarimax_holdout(feature_df, "seatemperature", reduced_cols, "ARIMA_C")

metrics_arimax = pd.DataFrame([
    {"model": arimax_a["label"], "aic": arimax_a["aic"], "rmse": arimax_a["rmse"]},
    {"model": arimax_b["label"], "aic": arimax_b["aic"], "rmse": arimax_b["rmse"]},
    {"model": arimax_c["label"], "aic": arimax_c["aic"], "rmse": arimax_c["rmse"]}
]).sort_values("rmse")

metrics_arimax

```

Out[9]:

	model	aic	rmse	mae
0	ARIMAX_A_base	91.542874	1.154612	1.066847
2	ARIMAX_C_reduced_minus_mean_relative_humidity_...	71.648238	2.196574	1.779561
1	ARIMAX_B_base_plus_lag	73.636473	2.229038	1.811506

In [10]: plt.figure(figsize=(12, 5))

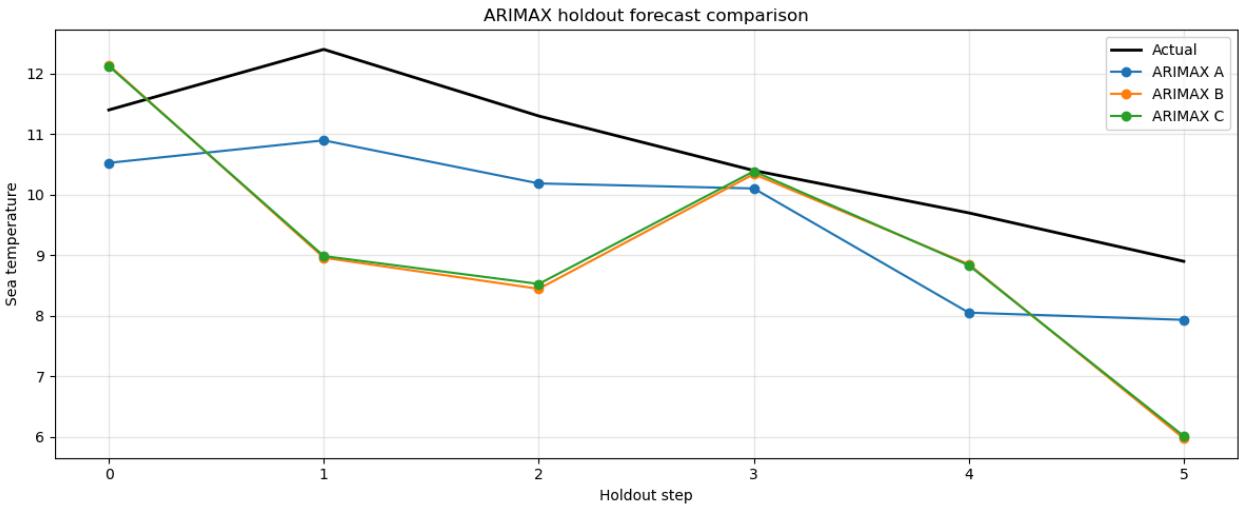
```

x_idx = np.arange(len(arimax_a["actual"]))
plt.plot(x_idx, arimax_a["actual"], color="black", linewidth=2, label="Actual")
plt.plot(x_idx, arimax_a["predicted"], marker="o", label="ARIMAX A")
plt.plot(x_idx, arimax_b["predicted"], marker="o", label="ARIMAX B")
plt.plot(x_idx, arimax_c["predicted"], marker="o", label="ARIMAX C")

plt.title("ARIMAX holdout forecast comparison")
plt.xlabel("Holdout step")
plt.ylabel("Sea temperature")
plt.grid(alpha=0.3)
plt.legend()
plt.tight_layout()

```

```
plt.show()
```



```
In [11]: ml_df = feature_df[["seatemperature"] + (base_cols + lag_cols)].dropna().reset_index()
X = ml_df[base_cols + lag_cols].astype(float)
y = ml_df["seatemperature"].astype(float)

split_idx = max(12, len(ml_df) - 6)
X_train, X_test = X.iloc[:split_idx], X.iloc[split_idx:]
y_train, y_test = y.iloc[:split_idx], y.iloc[split_idx:]

ts_cv = TimeSeriesSplit(n_splits=3)
param_grid = {
    "n_estimators": [100, 200, 300],
    "learning_rate": [0.03, 0.05, 0.1],
    "max_depth": [2, 3, 4],
    "subsample": [0.8, 1.0],
}

search = GridSearchCV(
    estimator=GradientBoostingRegressor(random_state=42),
    param_grid=param_grid,
    cv=ts_cv,
    scoring="neg_mean_squared_error",
    n_jobs=-1,
)
search.fit(X_train, y_train)

best_gb = search.best_estimator_
gb_pred = best_gb.predict(X_test)

gb_rmse = float(np.sqrt(mean_squared_error(y_test, gb_pred)))
gb_mae = float(mean_absolute_error(y_test, gb_pred))

metrics_ml = pd.DataFrame([
    {
        "model": "GradientBoosting_best_cv",
        "rmse": gb_rmse,
        "mae": gb_mae,
    }
])
```

```

        "best_params": str(search.best_params_),
    }
]

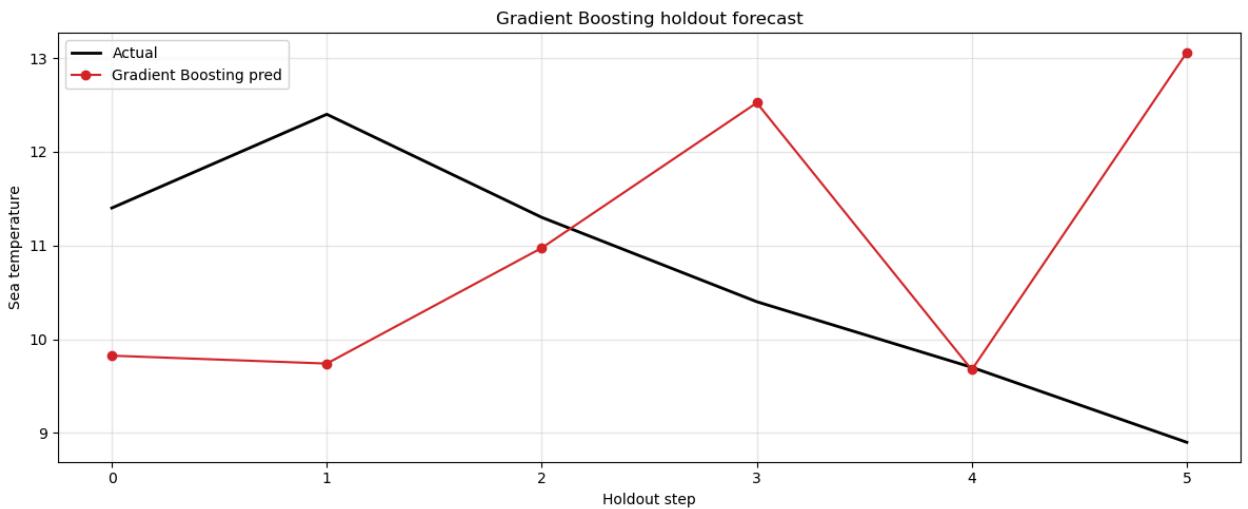
metrics_ml
```

Out[11]:

	model	rmse	mae	best_params
0	GradientBoosting_best_cv	2.290185	1.81152	{'learning_rate': 0.1, 'max_depth': 2, 'n_estimators': 100}

In [12]:

```
plt.figure(figsize=(12, 5))
plt.plot(np.arange(len(y_test)), y_test.to_numpy(), color="black", linewidth=2)
plt.plot(np.arange(len(gb_pred)), gb_pred, marker="o", color="tab:red", label="Gradient Boosting pred")
plt.title("Gradient Boosting holdout forecast")
plt.xlabel("Holdout step")
plt.ylabel("Sea temperature")
plt.grid(alpha=0.3)
plt.legend()
plt.tight_layout()
plt.show()
```



6) Save outputs

In [13]:

```
week_region_out = OUTPUT_DIR / "week_region_pivot.csv"
pd_ila_out = OUTPUT_DIR / "pd_ila_lat_pivot.csv"
feature_out = OUTPUT_DIR / "feature_df.csv"
sliding_out = OUTPUT_DIR / "sliding_window_best.csv"
arimax_metrics_out = OUTPUT_DIR / "arimax_metrics.csv"
ml_metrics_out = OUTPUT_DIR / "ml_metrics.csv"

pred_out = OUTPUT_DIR / "arimax_predictions_holdout.csv"
gb_pred_out = OUTPUT_DIR / "gb_predictions_holdout.csv"

pivot_table_week_region.reset_index().to_csv(week_region_out, index=False)
pivot_table_pd_ila.to_csv(pd_ila_out, index=False)
```

```

feature_df.to_csv(feature_out, index=False)
best_cfg_df.to_csv(sliding_out, index=False)
metrics_arimax.to_csv(arimax_metrics_out, index=False)
metrics_ml.to_csv(ml_metrics_out, index=False)

pd.DataFrame({
    "actual": arimax_a["actual"],
    "arimax_a_pred": arimax_a["predicted"],
    "arimax_b_pred": arimax_b["predicted"],
    "arimax_c_pred": arimax_c["predicted"],
}).to_csv(pred_out, index=False)

pd.DataFrame({
    "actual": y_test.to_numpy(),
    "gb_pred": gb_pred,
}).to_csv(gb_pred_out, index=False)

print(f"Saved: {week_region_out}")
print(f"Saved: {pd_ila_out}")
print(f"Saved: {feature_out}")
print(f"Saved: {sliding_out}")
print(f"Saved: {arimax_metrics_out}")
print(f"Saved: {ml_metrics_out}")
print(f"Saved: {pred_out}")
print(f"Saved: {gb_pred_out}")

```

Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/week_region_pivot.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/pd_ila_lat_pivot.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/feature_df.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/sliding_window_best.cs
v
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/arimax_metrics.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/ml_metrics.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/arimax_predictions_hol
dout.csv
Saved: /Users/endreasgard/NMBU/IND320/IND320/CA3/outputs/gb_predictions_holdou
t.csv

In [14]:

```

spark.stop()
cluster.shutdown()
print("Closed Spark and Cassandra connections.")

```

Closed Spark and Cassandra connections.